

Department of Information Technology and Electrical Engineering

Machine Learning on Microcontrollers

227-0155-00L

Exercise 6

Quantization and Pruning for Image Classification on STM32U5

ETH Center for Project-Based Learning

Wednesday 22nd October, 2025

1 Introduction

During the last exercise, you learned how to use the *TFLite* toolchain to execute TensorFlow Lite models on different Microcontroller (MCU) targets. We will come back to the same approach for the B-U585I-IOT02A board and perform real-time inference on the MCU. This is done using a simple *Python* script that establishes a Universal Asynchronous Receiver Transmitter (UART) connection to the MCU, sends the input data, and receives the classification result after inference. The data transfer will be done in a static fashion, meaning we will store the data locally and load it into the MCU from a *.npy*-array file. However, for your projects, you could also use the onboard sensors or other inputs such as your laptop camera, and transmit the data using a serial port to the MCU. To prepare you for your final assignment we will follow a bottom-up approach in creating a new project enabling real-time inference on your target architecture.

2 Notation

Student Task: Parts of the exercise that require you to complete a task will be explained in a shaded box like this.

Note: You find notes and remarks in boxes like this one.

3 Preparation

We will use the STM32 CUBE IDE which you are already familiar with in order to port the *TFLite* model on your microcontroller.

All Python dependencies required for feature extraction and training (e.g., *numpy*, *scipy*, *tensorflow*, etc.) are provided via a **Docker setup**.

Please make sure you have installed the environment in the first exercise on your system before you start.

Now, please activate your Docker environment and make your workspace ready.

Starting a new project

In earlier exercises, you learned how to create a new project in the STM32 CUBE IDE. Revise the previous materials if you do not remember how it is done. Carefully follow these steps and read the task descriptions thoroughly to avoid extra work.

Student Task 1 (Configuring a new project with UART):

1. Create a new project with STM32CubeMX.
2. Clear the default pinout and configure the USART1 peripheral on pins PA9 and PA10 as done in the previous exercises.
3. Switch to the Project Manager tab. Name the project and select STM32CubeIDE as the toolchain / IDE.
4. Switch to Code Generator and select to generate the peripheral initialization as a pair of '.c/.h' files per peripheral.
5. Generate Code. Save the project in the STM32 CUBE IDE workspace and import it in STM32 CUBE IDE .
6. In STM32 CUBE IDE , right-click the project and convert it to C++. We will need this, as the tensorflow lite library we will use is written in C++.




Figure 1: Project Manager configuration for separating the peripherals and main function.




Figure 2: After generating the code, import the project to STM32 CUBE IDE .




Figure 3: Convert the project to C++ using STM32 CUBE IDE .

Compiling and flashing a project

In the previous exercises, it was stated that the project can be compiled by *simply clicking the BUILD button* as shown in Figure 4.

However, this is not entirely true as you have to be very careful when starting your own exploration.

We will give you a basic intuition of what can go wrong and how to fix it. In Figure 5, a general compilation flow is shown. We start with a simple C program. C is a *compiled language*, i.e. a compiler is needed to translate source code to machine-readable code. One of the most famous ones is the GNU Compiler Collection (GCC), which is also used in the STM32 CUBE IDE. Compilers translate source code in four steps: 1) Preprocessing, 2) Compiling, 3) Assembling, and 4) Linking. If you want to read up on the different steps you can refer to this blog post.

We will focus mainly on the **Linking** part, as this is usually where things go wrong. The **Assembler** goes through all the files in your C/C++ project and generates so-called *object code (binary)*. This is essentially pure machine code, which runs on your target architecture. However, since we have multiple files in our project such as header files, or external libraries, we need the **Linker** to combine everything into a single *executable*. As the name suggests, this file contains the code which is in the end run on the MCU to perform the inference.

However, the **Linker** needs some help when using libraries and external files such as *CMSIS* or *TFLite*. Both tools contain a hierarchy of definitions and macros. Macros are essentially function definitions used across several source files. These definitions are usually all listed in so-called header files. You request the use of a header file in your program by *including* it, with the C preprocessing directive `#include`. Thus, you can access the macros within the *included* header file.

In Figure 6, an example of hierarchical header files is shown, as they occur e.g. in the *CMSIS* library. *header_three.h* has the lowest hierarchy and includes macros from the other two header files. In order for a project to successfully compile we have to tell the **Linker** the order of macro definitions. Otherwise, it will try to link functions that have not yet been defined. If you encounter error messages during building such `undefined reference to XYZ` it is probably due to misordering of the paths within your project. You can check the order of the search paths under `Properties→Paths` and `\ Symbols→Includes`.




Figure 4: The Build and Debug icons in the Toolbar.




Figure 5: General compilation flow.




Figure 6: Header file structure example.

Student Task 2 (Project preparation):

1. Download the exercise materials from the course page.
2. Copy the `tensorflow_lite` folder to the projects root.
3. Copy `CMSIS/DSP`, `CMSIS/NN` and `CMSIS/Core` to the existing `Drivers/CMSIS` folder.
4. Under `Core→Inc` add a new folder called `models`. This is where we will store the TFLite model header files for the inference step on the MCU.

Note: When adding the includes, do this through the IDE's file browser. Also check that you have a Symbol `CMSIS_NN` and make sure to add `tensorflow_lite` as a source location in the `Paths` and `Symbols` section of the Project properties.

Note: In this exercise, we have provided the *CMSIS* and *TFLite* toolchains for you. If you want to update the toolchains for your own projects you can find both *CMSIS* and *TFLite* on GitHub. It is usually good to check these repositories from time to time, as more and more kernels are added to these toolchains, which could improve your performance and accuracy significantly.

We will now add the folders to the include path. This can be tricky to get right for compilation to work properly. To add or remove a folder to the include path, right-click the folder in the Project Explorer and click

Add/remove include path... as in figure 7.




Figure 7: How to add a folder to the include path.

Student Task 3 (Include Directory Management):

1. Add the folders shown in figure 8 to the include path. Add them in this exact order.
2. Open the window:
Project / Properties / C/C++ General / Paths and Symbols
and confirm the added includes match figure 8 for both GNU C and GNU C++ languages.
3. Switch to the Symbols tab and add a symbol CMSIS_NN without a value.
Make sure to select both:
Add to all configurations
Add to all languages
4. Switch to the Source Location tab. Add the folder /tensorflow_lite.

 <code>/\${ProjName}/Drivers/CMSIS/Core/Include</code>
 <code>/\${ProjName}/Drivers/CMSIS/DSP/Include</code>
 <code>/\${ProjName}/Drivers/CMSIS/DSP/PrivateInclude</code>
 <code>/\${ProjName}/Drivers/CMSIS/NN/Include</code>
 <code>/\${ProjName}/tensorflow_lite</code>
 <code>/\${ProjName}/tensorflow_lite/third_party/flatbuffers/include</code>
 <code>/\${ProjName}/tensorflow_lite/third_party/gemmlowp</code>
 <code>/\${ProjName}/tensorflow_lite/third_party/kissfft</code>
 <code>/\${ProjName}/tensorflow_lite/third_party/ruy</code>
 <code>/\${ProjName}/Core/Inc/models</code>

Figure 8: The necessary include directories.

3.1 Printing output with UART

Since we want to be able to communicate via UART we will add a function to write characters to the serial communication stream. Thus, we can communicate out of the microcontroller via UART. This can be useful for debugging purposes. **Do not** flash the MCU yet.

Student Task 4:

1. Revise the previous exercise material on how to include the `stdio.h` header.
2. Add a function definition for `_write` to write to the output stream in the appropriate code section of the main file.
3. Add a `printf` statement to the main function body to announce the beginning of the inference; this will furthermore test the serial communication. You can remove this for the live inference test later on.

4 Preparing the Inference step on the MCU

In the exercise materials we provide you with a `Jupyter notebook` to convert a trained *Keras* model via *TFLite* to a *C header file*. The header file contains all the network parameter information used for *inference*. We will start with a very simple dataset, i.e. the Modified National Institute of Standards and Technology (MNIST) dataset and a straightforward network implementation. Afterwards, you will implement your own network, targeting classification on the CIFAR-10 dataset.

Student Task 5:

1. Open the `Jupyter notebook` file `lab7.ipynb` inside your Docker, using JupyterLab. We will start with **Task 1**.
2. What is the validation loss and accuracy that you can achieve on the test set?
 - Validation loss:

Solution: It should be around 0.29.

- Validation accuracy:

Solution: Around 89%.

3. Briefly explain the difference between the test and validation set.

- Validation set:

Solution: The validation set is used **during** training to tune the hyperparameters of the model to achieve best performance.

- Test set:

Solution: The test set is used to evaluate the model's final performance on data it has not seen yet.

4. What is the size of the `.h5` model and the `.tflite` model? By how much can we reduce the model size?

- Size of the `.h5` file:

Solution: 264 KB

- Size of the `.tflite` file:

Solution: 82KB

- Compression factor:

Solution: $3.21\times$

Now you have successfully converted your *Keras* model to a *TFLite* model for the MCU.

5 Deploying the network to the MCU for inference

Deploying the network effectively means storing the network's parameters (e.g., weights and biases), usually in the MCU's read-only memory, as well as generating the C code implementing the network's computational graph, managing intermediate buffer memory and calling (optimized) kernel implementations of individual layers. Therefore, the first prerequisite for a successful deployment is ensuring that the network's size, given the number of parameters and their precision, is smaller than the available storage space.

Student Task 6:

1. Move on to **Task 2** in the `Jupyter Notebook`.
2. Count the model parameters, considering the weights and biases within each layer. Verify your results by comparing them with the output of the `summary()`^a method.

Solution: The convolutional layer produces $c = (28 - 3 + 1)^2 \times 12$ activations. This results in $m = (c/2)^2 \times 12$ activations after performing maxpooling. Finally, we have 10 activations in the dense layer. This results in a total number of parameters:

$$\#num_params = (m + 1) \times 10 + c = 20410$$

3. By how much can you reduce the model size by performing full 8-bit quantization?

Solution: The reduction factor is $3.48\times$.

4. What accuracy can we achieve with the fully quantized model? Why do you think the quantized model might achieve higher accuracy than the full precision model?

Solution: The PTQ accuracy is 89.52%. It is slightly higher, as quantization can serve as a form of dropout mitigating overfitting of the model.

^a Note that similar functionality is covered in PyTorch by the `torchsummary` library.

Another hardware-associated constraint that must be addressed when deploying a model on an MCU is represented by the memory limitations. These limitations refer to the available read-write memory used for the intermediate buffers; in the absence of tiling and multi-buffered memory accesses, the memory requirements of a network can be defined as:

$$M = \max_{l \in L} l_{in} + l_{out} + l_p \quad (1)$$

, where l_{in} represent the input activations of a layer l of a network comprised of L layers, l_{out} represent the output activations, whilst l_p are the layer's parameters. Similarly to the storage requirements, the memory limitations also depend on the precision used to represent the data.

Student Task 7:

1. What are the memory requirements of your network? Do they fit the constraints of your target platform?

Solution: There are 20410 model parameters. For the convolutional layer there are 28×28 input and $26 \times 26 \times 12$ output activations. These output activations are also the input activations of the maxpooling layer. The number of output activations here is $13 \times 13 \times 12$. The dense layer has only 10 output activations. This results in a total of 31344 values that must be stored. For the FP32 model this means 123 KB and for the INT8 model 31KB.

In the next step, we have to include the trained network in our project and deploy it on the MCU.

Student Task 8:

1. In the exercise material we provide you the application source code and header file to run the inference on the MCU. Add the `.h` and `.cpp` files at the right locations to your project.
2. Add an `include` directive for the application `.h` file in the code and invoke the `application` function after all initialization code.
3. Add the model `.h` file that you generated from the `Jupyter notebook` and the initial data `.h` to your model directory in your project. Make sure you include the correct model in the application code.

4. Now you can compile and flash your project. Once you verify that the build is successful and the application is started, the `printf` statement from the main function body, announcing the inference step, can be commented out.

Note: To measure the inference duration of your model use the *CycleCounter* as before. To use it in `app.cpp` you have to include it using `extern "C" {#include "CycleCounter.h"}`. This will ensure that the functions declared in `CycleCounter.h` are treated correctly by both C and C++ compilers, resolving the conflicting declaration error.

Note: Our real-time inference script in the next tasks reads the bytes directly from the serial stream. Make sure that you do not have any `printf` statements such as e.g. from the *CycleCounter* in your code that contaminate the stream.

6 Real-time inference on the MCU

Finally, we can perform the inference on the MCU. For this purpose, we provide you with a small Graphical User Interface (GUI), which is programmed in the `test.py` script provided in the exercise materials. In Figure 9, the general working principle of the test script is shown. We send the test image data together with its label to the MCU. After one inference step we read out the UART port to retrieve the predicted label.

In order to achieve real-time operation for our system, our model's latency (i.e., the interval between the model receiving the input and said model producing the prediction) has to be smaller than the data acquisition time, the latter being emulated here using the `test.py` script. The inference latency can thus be considered a third hardware-associated constraint, further determining the energy consumption (i.e., $E = P \cdot t$) of our system. Although the number of Floating Point Operations (FLOPs) might represent a sufficiently good proxy for the latency when comparing different networks with a similar architecture, optimizations such as the usage of Single Instruction Multiple Data (SIMD)-based kernels, tiling, or double buffering could make the FLOPs-based comparison obsolete. It is thus recommended, when optimizing a neural network considering an accuracy-latency trade-off, to perform hardware-in-the-loop optimizations by measuring the network's latency on the target platform.




Figure 9: Real-time Inference step on the MCU.

Student Task 9:

1. Navigate to the location of your `inference.py` script.
2. Check out the source code of the `inference.py` script and try to understand how it works.
3. To perform the inference run the following command:
`python inference.py <dataset_name>`. *Attention:* You might have to modify the script slightly, for your communication interface.
4. What is the latency and memory usage of your network?

Solution: The model execution latency should be around 9M cycles (about 17.8 frames/second @ 160MHz). The overall RAM utilization is about 90KB.

Congratulations, you managed to run real-time inference on the microcontroller! As you have probably seen, the `inference.py` script supports also other datasets, such as *Fashion MNIST* and *CIFAR-10*.

7 Advanced Quantization Techniques

So far, we have seen *Post-training quantization* (PTQ). However, there exist more elaborate quantization techniques, such as *Quantization-aware training* (QAT). For our simple model we do not observe any loss in accuracy, but in more complex models, the accuracy loss can be significant. To accommodate for this, we can use QAT instead. In QAT, we train the model with quantization already in mind. The drawback of this approach is that we need more time to train the model. Depending on your application, different techniques can be more suitable. In Figure 10 we provide you with a decision tree to help you choose the right quantization technique.




Figure 10: Decision Tree for choosing the right quantization technique. Source: TensorFlow

Student Task 10:

1. Move on to **Task 3** in the `Jupyter` notebook.
2. What are the accuracy and loss we can achieve on the test set with QAT?
 - Accuracy:
Solution: Around 90%.
 - Loss:
Solution: Around 0.28.
3. Do we sacrifice memory for the accuracy gain? Do you think it is worth it?

Solution: We only have a slight increase in memory due to the additional trainable parameters in QAT, however the accuracy improves which justifies the overhead.

8 Pruning

In the previous tasks, we have seen how to quantize a model. However, we can also reduce the number of parameters in a model. This is referred to as *pruning*. Pruning in machine learning refers to the technique of reducing the size and complexity of a trained model by removing unnecessary or redundant parameters or connections. By redundant, we try to remove parameters that do not contribute to the model's accuracy (i.e. non-critical weights). The goal of pruning is to make the model

more efficient in terms of memory usage, computation speed, and energy consumption, without sacrificing its performance on the task it was trained for.

Student Task 11:

1. Move on to **Task 4** in the `Jupyter notebook`.
2. What are the accuracy and loss we can achieve on the test set with the structured pruned model using constant sparsity?

- Accuracy:

Solution: Around 90%.

- Loss:

Solution: Around 0.28.

3. What is the compression factor we can achieve with structured pruning using constant sparsity?

Solution: We can achieve $1.62\times$ compression.

4. Write down the results for unstructured pruning with constant sparsity.

- Accuracy:

Solution: Around 90%.

- Loss:

Solution: Around 0.28.

- Compression factor:

Solution: The compression ratio is $1.61\times$.

5. Is the accuracy for the unstructured or structured pruning better? How can you explain your observations?

Solution: They are almost the same. This is due to the fact that the model is not extremely complex. In general, structured pruning usually achieves higher accuracy, because the model shape remains the same, however, unstructured pruning is more effective w.r.t. reducing the model's size.

6. Write down the results for unstructured pruning with dynamic sparsity.

- Accuracy:

Solution: Around 90%.

- Loss:

Solution: Around 0.29.

- Compression factor:

Solution: We can achieve a compression of $1.02\times$.

7. Why do you think our compression factor is now lower compared to constant sparsity pruning?

Solution: With the dynamic schedule we account for more weights still being present in the network, as this sparsity scheme is less aggressive compared to constant sparsity.

9 Pruning-preserving Quantization-aware Training

We saw that we can reduce the model size by means of pruning, however, we sacrifice accuracy. In order to mitigate the accuracy loss, we can combine pruning with quantization-aware training. Thus, we can find a good trade-off between accuracy and model size.

Student Task 12:

1. Move on to **Task 5** in the `Jupyter notebook`.
2. What are the accuracy and loss we can achieve on the test set with pruning-preserving QAT?
 - Accuracy:
Solution: Around 90%.
 - Loss:
Solution: Around 0.28.
 - Compression factor:
Solution: The compression factor is $4.61\times$
3. Try to interpret the table and plot in the *Result Summary* section of the `Jupyter notebook`.

Solution: We see that pruning-preserving outperforms all the other methods in terms of accuracy and model size. Interestingly, the full-precision model performs the worst. This is probably due to the fact that the model is very simple and overfitting, which can be mitigated by quantization. The pruned models without QAT show significant memory overhead compared to PTQ, which has to be considered depending on your target application. On the other hand, all the pruned model increase the model's accuracy.

Now you have generated all the files necessary to run the inference of each model on the microcontroller.

Student Task 13:

1. Repeat the steps for running the inference as described in Student Task 9.
2. What is the latency of each model?
 - *Post-training quantization*:
Solution: 14M Cycles
 - *Quantization-aware training*:

Solution: 14M Cycles

- *Constant Structured Pruning:*

Solution: 46M Cycles

- *Constant Unstructured Pruning:*

Solution: 46M Cycles

- *Dynamic Unstructured Pruning:*

Solution: 46M Cycles

- *Pruning-preserving quantization-aware training:*

Solution: 14M Cycles

3. What is the memory usage of each model?

- *Post-training quantization:*

Solution: 571 kB

- *Quantization-aware training:*

Solution: 572 kB

- *Constant Structured Pruning:*

Solution: 630 kB

- *Constant Unstructured Pruning:*

Solution: 630 kB

- *Dynamic Unstructured Pruning:*

Solution: 630 kB

- *Pruning-preserving quantization-aware training:*

Solution: 572 kB



Congratulations! You have reached the end of the exercise.
If you are unsure of your results, discuss them with an assistant.

